

TDDE07 - Lab 3

Pontus Svensson (ponsv690) & Kristian Sikiric (krisi211)


Assignment 1


Normal model, mixture of normal model with semi-conjugate prior

In this assignment a data file containing the daily precepitation from year 1948 to 1983 was given. We were to analyze this data using two different models using gibbs sampling. The daily precipitation y_1, \dots, y_n were assumed to be independently normal distributed, $y_1, \dots, y_n | \mu, \sigma^2 \sim \mathcal{N}(\mu, \sigma^2)$ were $\mu \sim \mathcal{N}(\mu_0, \tau_0^2)$ and $\sigma \sim \text{Inv} - \mathcal{X}^2(\nu_0, \sigma_0^2)$. Since there are two random variables, we have a bivariate normal distribution.


The full conditional posteriors were given to be $\mu | \sigma^2, x \sim \mathcal{N}(\mu_n, \tau_n^2)$ and $\sigma^2 | \mu, x \sim \text{Inv} - \mathcal{X}(\nu_n, \frac{\nu_0 \sigma_0^2 + \sum_{i=1}^n (x_i - \mu)^2}{n + \nu_0})$


Now we could use these to implement a gibbs sampler, the initial variables were set to some random values based on our priors, (see code). The following histogram shows the result after sampling.



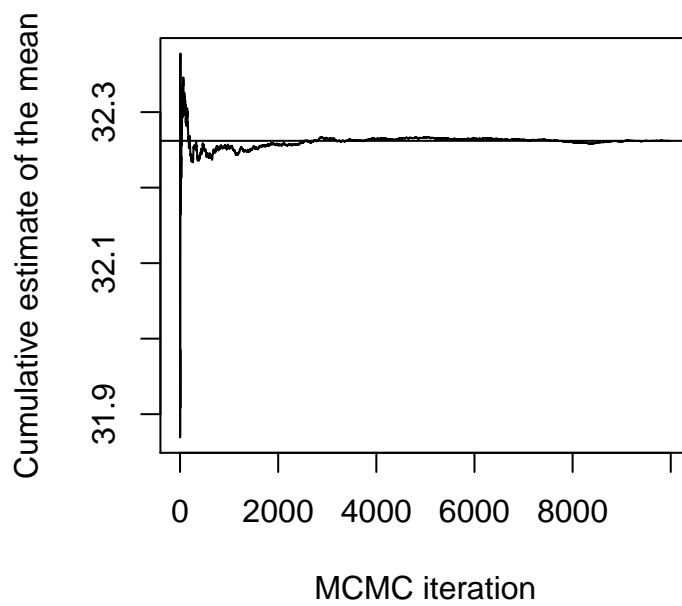



These plots says that the mean precipitation in this period was around $\frac{32.5}{100}$ inches with a variance of $\frac{40}{100}$ inches. The following plots shows the convergence of the sampler.





Cusum – Gibbs





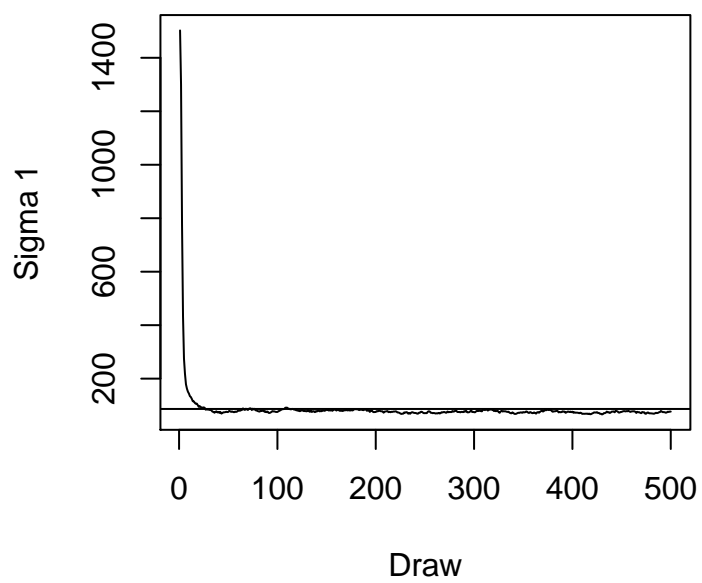




The plots show that the sampler converges. The first two plots of the trajectories shows that the different samples are not so correlated, which is good.

We then repeated the process but now the data followed a two-component mixture of normals model. The code was given, we only extended it so we could show the convergence of the sampler. Below are some plots showing the results from the sampler with the convergence for the different parameters. We can see that perhaps the first 100 iterations can be seen as the burn-in period.

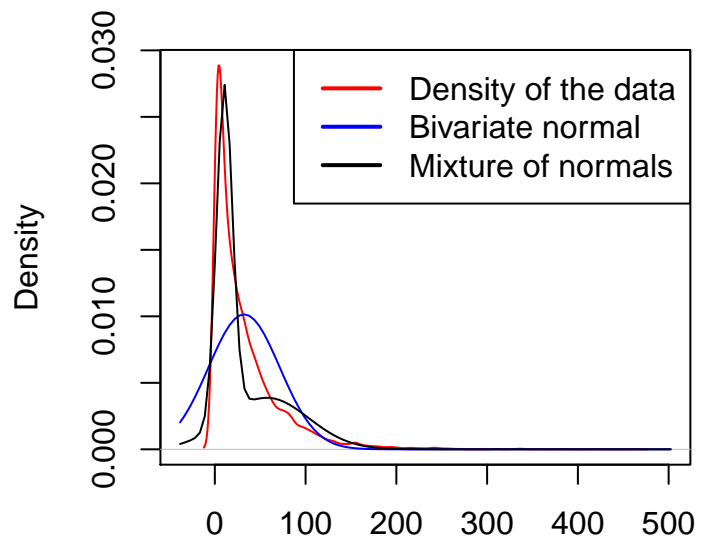
Final fitted density







Below is a graphical comparison between the density of the data, the sampled normal density from above aswell as the sampled mixture of normals.



Assignment 2

Metropolis Random Walk for Poisson regression

In this assignment we considered the following poisson regression model $y_i|\beta \sim \text{Poisson}[\exp(x_i^T \beta)]$, $i = 1, \dots, n$ where y_i is the count for the i th observation in the sample and x_p is the p -dimensional vector with covariates. A data set containing observations from 1000 eBay auctions of coins were given. The target variable in this assignment was the number of bids in each auction.

First a general linear model was fitted using the `glm` function in R. From this model we could see that the covariates Sealed, VerifyId and MinBidShare were the most significant.

Now we did a Bayesian analysis of the Poisson regression. The prior was $\beta \sim \mathcal{N}[0, 100 * (X^T X)^{-1}]$ where X is the $n \times p$ covariate matrix. The posterior density was assumed to be approximately multivariate normal: $\beta|y \sim \mathcal{N}(\tilde{\beta}, J_y^{-1}(\tilde{\beta}))$. $\tilde{\beta}$ and $J_y(\tilde{\beta})$ were computed with the `optim` function in R with the log posterior function of the Poisson model. Below the coefficients from the Bayesian analysis with `optim` is compared to the coefficients of the `glm` model, as we can see, they are very similar.



Bayesian analysis



```
## [1]  1.07235230 -0.02033338 -0.39274453  0.44366399 -0.05220897 -0.22145707
## [7]  0.07033797 -0.12041259 -1.89392696
```



GLM model



```
##      Const PowerSeller   VerifyID      Sealed      Minblem      MajBlem
##  1.07244206 -0.02054076 -0.39451647  0.44384257 -0.05219829 -0.22087119
##      LargNeg      LogBook MinBidShare
##  0.07067246 -0.12067761 -1.89409664
```

Finally, we simulated from the actual posterior of β using the Metropolis algorithm. Our proposal density was $\theta_p|\theta^{(i-1)} \sim \mathcal{N}(\theta^{(i-1)}, c \sum)$, where $\sum = J_y^{-1}(\tilde{\beta})$ was obtained from the `optim` function. The value c is a tuning parameter, we used $c = 2$ in this assignment. When simulating with the metropolis algorithm using the same log posterior mentioned above, we got the following convergence curves for the different covariates. The samples does seem a bit correlated, but the result seems good still.









Here we also notice that perhaps the first 1000 iterations can be seen as the burn-in.

And we got the following values for the coefficients, which does not differ much from when using `optim` in the bayesian analysys.

Metropolis algorithm

```
## [1]  1.06104095 -0.01993666 -0.39559328  0.43507861 -0.05283118 -0.23656841
## [7]  0.07295384 -0.12145117 -1.88318986
```

Appendix

Assignment 1

```
data = read.delim("/Users/kristiansikiric/Desktop/TDDE07/Lab3/rainfall.dat")
#data = read.delim("/home/krisi211/Desktop/TDDE07/Lab3/rainfall.dat")
set.seed(123)
## a)
# Init values
x = mean(data[,1])
n = length(data[,1])

##### Random values #####
mu_0 = x
tau_0 = 10
nu_0 = 3
sigma_0 = 1
sigma = 1 #Init sigma to some value not zero
#####

nu_n = nu_0 + n
NDraws = 10000

#Gibbs sampling
gibbsDraws = matrix(0,NDraws,2)
for( i in 1:NDraws){

  #####FROM LECTURE 2#####
  w = (n/sigma) / (n/sigma + 1/tau_0)
  mu_n = w*x + (1-w)*mu_0
  tau_n = 1/((n/sigma) + (1/tau_0))
  #####

  mu_gibbs = rnorm(1,mu_n,sqrt(tau_n))
  gibbsDraws[i,1] = mu_gibbs

  tau = (nu_0*sigma_0 + sum((data[,1]-mu_gibbs)^2))/(n+nu_0)
  sigma = ((nu_n-1)*tau)/rchisq(1,nu_n-1)
  gibbsDraws[i,2] = sigma
}

hist(gibbsDraws[,1])
hist(gibbsDraws[,2])

plot(gibbsDraws[,1],type = 'l')
plot(gibbsDraws[,2],type = 'l')
```

```

cusumData = cumsum(gibbsDraws[,1])/seq(1,NDraws)
plot(1:NDraws, cusumData, type = "l", ylab='Cumulative estimate',
     xlab = 'MCMC iteration',
     xlim = c(0,NDraws),
     main = 'Cusum - Gibbs')
abline(h = mean(gibbsDraws[,1]))

cusumData = cumsum(gibbsDraws[,2])/seq(1,NDraws)
plot(1:NDraws, cusumData, type = "l", ylab='Cumulative estimate',
     xlab = 'MCMC iteration',
     xlim = c(0,NDraws),
     main = 'Cusum - Gibbs')
abline(h = mean(gibbsDraws[,2]))

## b)
#source("/home/krisi211/Desktop/TDDE07/Lab3/NormalMixtureGibbs.R")
source("/Users/kristiansikiric/Desktop/TDDE07/Lab3/NormalMixtureGibbs.R")

## c)
plot(density(data[,1]),col='red',xlim = c(xGridMin,xGridMax))
mu_hat = mean(gibbsDraws[,1])
sigma_hat = mean(gibbsDraws[,2])
lines(xGrid,dnorm(xGrid,mu_hat,sqrt(sigma_hat)), col = 'blue')
lines(xGrid,mixDensMean, type = "l", lwd = 2, lty = 4, col = "black")

```

NormalMixtureGibbs.r

```

# Estimating a simple mixture of normals
# Author: Mattias Villani, IDA, Linköping University. http://mattiasvillani.com

##### BEGIN USER INPUT #####
# Data options
rawData <- read.delim("/Users/kristiansikiric/Desktop/TDDE07/Lab3/rainfall.dat")
x <- as.matrix(rawData['X136'])

# Model options
nComp <- 2 # Number of mixture components

# Prior options
alpha <- 10*rep(1,nComp) # Dirichlet(alpha)
muPrior <- rep(mean(x),nComp) # Prior mean of mu
tau2Prior <- rep(10,nComp) # Prior std of mu
sigma2_0 <- rep(var(x),nComp) # s20 (best guess of sigma2)
nu0 <- rep(2,nComp) # degrees of freedom for prior on sigma2

# MCMC options
nIter <- 500 # Number of Gibbs sampling draws

# Plotting options
plotFit <- TRUE
lineColors <- c("blue", "green", "magenta", "yellow")
sleepTime <- 0.1 # Adding sleep time between iterations for plotting
##### END USER INPUT #####

```

```
##### Defining a function that simulates from the
rScaledInvChi2 <- function(n, df, scale){
  return((df*scale)/rchisq(n,df=df))
}

##### Defining a function that simulates from a Dirichlet distribution
rDirichlet <- function(param){
  nCat <- length(param)
  piDraws <- matrix(NA,nCat,1)
  for (j in 1:nCat){
    piDraws[j] <- rgamma(1,param[j],1)
  }
  # Dividing every column of piDraws by the sum of the elements in that column.
  piDraws = piDraws/sum(piDraws)
  return(piDraws)
}

# Simple function that converts between two different representations of the mixture allocation
S2alloc <- function(S){
  n <- dim(S)[1]
  alloc <- rep(0,n)
  for (i in 1:n){
    alloc[i] <- which(S[i,] == 1)
  }
  return(alloc)
}

# Initial value for the MCMC
nObs <- length(x)
S <- t(rmultinom(nObs, size = 1 , prob = rep(1/nComp,nComp)))
mu <- quantile(x, probs = seq(0,1,length = nComp))
sigma2 <- rep(var(x),nComp)
probObsInComp <- rep(NA, nComp)

# Setting up the plot
xGrid <- seq(min(x)-1*apply(x,2,sd),max(x)+1*apply(x,2,sd),length = 100)
xGridMin <- min(xGrid)
xGridMax <- max(xGrid)
mixDensMean <- rep(0,length(xGrid))
effIterCount <- 0
ylim <- c(0,2*max(hist(x,plot=FALSE)$density))
print = TRUE

mixDraws = matrix(0,nIter,4)

for (k in 1:nIter){
  #message(paste('Iteration number:',k))
  alloc <- S2alloc(S)
  nAlloc <- colSums(S)
  #print(nAlloc)
  # Update components probabilities
  pi <- rDirichlet(alpha + nAlloc)
```

```

# Update mu's
for (j in 1:nComp){
  precPrior <- 1/tau2Prior[j]
  precData <- nAlloc[j]/sigma2[j]
  precPost <- precPrior + precData
  wPrior <- precPrior/precPost
  muPost <- wPrior*muPrior + (1-wPrior)*mean(x[alloc == j])
  tau2Post <- 1/precPost
  mu[j] <- rnorm(1, mean = muPost, sd = sqrt(tau2Post))
  mixDraws[k,j] = mu[j]
}

# Update sigma2's
for (j in 1:nComp){
  sigma2[j] <- rScaledInvChi2(1, df = nu0[j] + nAlloc[j],
                             scale = (nu0[j]*sigma2_0[j] +
                                       sum((x[alloc == j] -
                                             mu[j])^2))/(nu0[j] + nAlloc[j]))
  mixDraws[k,j+2] = sigma2[j]
}

# Update allocation
for (i in 1:nObs){
  for (j in 1:nComp){
    probObsInComp[j] <- pi[j]*dnorm(x[i], mean = mu[j], sd = sqrt(sigma2[j]))
  }
  S[i,] <- t(rmultinom(1, size = 1, prob = probObsInComp/sum(probObsInComp)))
}

# Printing the fitted density against data histogram
if (plotFit && (k%1 == 0) && print){
  effIterCount <- effIterCount + 1
  #hist(x, breaks = 20, freq = FALSE, xlim = c(xGridMin,xGridMax),
  #main = paste("Iteration number",k), ylim = ylim)
  mixDens <- rep(0,length(xGrid))
  components <- c()
  for (j in 1:nComp){
    compDens <- dnorm(xGrid,mu[j],sd = sqrt(sigma2[j]))
    mixDens <- mixDens + pi[j]*compDens
    #lines(xGrid, compDens, type = "l", lwd = 2, col = lineColors[j])
    components[j] <- paste("Component ",j)
  }
  mixDensMean <- ((effIterCount-1)*mixDensMean + mixDens)/effIterCount

  #lines(xGrid, mixDens, type = "l", lty = 2, lwd = 3, col = 'red')
  #legend("topleft", box.lty = 1, legend = c("Data histogram",components, 'Mixture'),
  #col = c("black",lineColors[1:nComp], 'red'), lwd = 2)
  #Sys.sleep(sleepTime)
}
}

```



```

hist(x, breaks = 20, freq = FALSE, xlim = c(xGridMin,xGridMax),
     main = "Final fitted density")
lines(xGrid, mixDensMean, type = "l", lwd = 2,
      lty = 4, col = "red")
lines(xGrid, dnorm(xGrid, mean = mean(x),
                  sd = apply(x,2,sd)),
      type = "l", lwd = 2, col = "blue")
legend("topright", box.lty = 1,
      legend = c("Data histogram","Mixture density","Normal density"),
      col=c("black","red","blue"), lwd = 2)

##### Our Code #####
plot(mixDraws[,1],type='l')
abline(h = mean(mixDraws[,1]))
plot(mixDraws[,2],type='l')
abline(h = mean(mixDraws[,2]))
plot(mixDraws[,3],type='l')
abline(h = mean(mixDraws[,3]))
plot(mixDraws[,4],type='l')
abline(h = mean(mixDraws[,4]))

```

Assignment 2

```

#data = read.delim("//Users/kristiansikiric/Desktop/TDDE07/Lab3/eBayNumberOfBidderData.dat", sep = "")
data = read.delim("/home/ponsv690/Documents/TDDE07/Lab3/eBayNumberOfBidderData.dat",sep = "")
set.seed(123)
X = as.matrix(data[,-1])
y = data[1]
## a)
glm.model= glm(nBids ~0+.,data = data, family = poisson)

#Significant covariates: Sealed, VerifyId, MajBlem(Semi), MinBidShare

## b)

mu = matrix(0,dim(X)[2],1)
sigma = 100 * solve(t(X)%*%X)
initVal = rep(0,dim(X)[2])
library("mvtnorm")

logPoisson = function(betas, y,X,mu,sigma){
  logPos = (sum(y*betas%*%t(X) - exp(betas%*%t(X)) - log(factorial(y))))

  if (abs(logPos) == Inf) logPos = -20001
  logPrior = dmvnorm(betas, mu, sigma)
  return(logPos + logPrior)
}

OptimResults = optim(initVal,logPoisson,gr=NULL,y,X,mu,
                    sigma,method=c("BFGS"),control=list(fnscale=-1),hessian=TRUE)
beta.tilde = OptimResults$par
inv.hessian = -solve(OptimResults$hessian)

```

```

beta = rmvnorm(10000, beta.tilde, inv.hessian)
(colMeans(beta))
coef(glm.model)

## c)
rwm = function(var, LogPost, theta_prev, ...){
  # Step 1
  theta_p = rmvnorm(1, theta_prev, var)

  # Step 2
  post.theta_prev = LogPost(theta_prev, ...)
  post.theta_p = LogPost(theta_p, ...)
  alpha = min(1, exp(post.theta_p - post.theta_prev))

  # Step 3
  accepted = runif(1, 0, 1) < alpha

  if(accepted){
    return(list(theta=theta_p, accepted=accepted))
  }
  else {
    return(list(theta=theta_prev, accepted=accepted))
  }
}

mcmc = function(LogPost, ndraws, ncov) {
  c = 2
  var = c * inv.hessian
  beta = matrix(rep(rep(0, ncov), ndraws), ncol = ncov)
  accepted = rep(0, ndraws-1)

  for(i in 2:ndraws) {
    sample = rwm(var, LogPost, beta[i-1,], y, X, mu, sigma)
    beta[i,] = sample$theta
    accepted[i-1] = sample$accepted
  }

  sum(accepted) / (ndraws-1)
  plot(beta[,2], type = 'l')
  plot(beta[,3], type = 'l')
  plot(beta[,4], type = 'l')
  plot(beta[,5], type = 'l')
  plot(beta[,6], type = 'l')
  plot(beta[,7], type = 'l')
  plot(beta[,8], type = 'l')
  plot(beta[,9], type = 'l')

  return(beta)
}

betas = mcmc(logPoisson, 10000, dim(X)[2])
colMeans(betas)
colMeans(beta) #From b)

```